

# Game-theoretic Foundations of Multi-agent Systems

## Lecture 7: Stochastic Games

Seyed Majid Zahedi

UNIVERSITY OF  
**WATERLOO**



# Outline

1. Markov Decision Processes
2. Definition of Stochastic Games
3. Strategies and Equilibria in Stochastic Games




## Motivation: Non-deterministic Search in Grid World




# Grid World Actions

## Deterministic




## Stochastic



# A Grid World Instance

- Agent lives in a grid
- Walls block agent's path
- Actions do not always go as planned
  - 80% of time, action "North" takes us north
  - 10% of time, "North" takes us west; 10% east
  - If there is a wall in the direction the agent would have been taken, the agent stays put
- Agent receives rewards at each step
- Goal is to maximize sum of rewards




# Change in Notation

- So far, we have used  $s$  to denote **strategy profile**
- In this lecture, we use  $\pi$  for strategy
- We use  $s$  to denote **state**



# Markov Property

- Given present state, future and past are independent
- Future state depends only on current state and action



Andrey Markov (1856-1922)

$$P(S_{t+1} = s \mid S_t = s_t, A_t = a_t, S_{t-1} = s_{t-1}, A_{t-1} = a_{t-1}, \dots, S_0 = s_0) = \\ P(S_{t+1} = s \mid S_t = s_t, A_t = a_t)$$



# Markov Decision Processes: Formal Definition


- $S$  is set of states and  $A$  is set of actions
- $p : S \times A \times S \mapsto [0, 1]$  specifies transition probabilities
  - $p(s, a, s')$  is probability of going to  $s'$  when taking  $a$  in  $s$
- $r : S \times A \times S \mapsto \mathbb{R}$  returns reward
  - $r(s, a, s')$  is reward of going to  $s'$  when taking  $a$  in  $s$
- There are two ways to aggregate rewards
  - Limit-average reward:  $\lim_{T \rightarrow \infty} \sum_{t=1}^T r_t / T$
  - **Future-discounted reward:**  $\sum_{t=1}^{\infty} \delta^{t-1} r_t$   
(we will focus on this)






## Example: Racing

- Autonomous car wants to travel far, quickly
- There are 3 states: Cool, Warm, Overheated, and 2 actions: **Slow**, **Fast**
- Going faster gets double reward



# Policies

- A (stationary, deterministic) policy  $\pi : S \mapsto A$  gives action for each state
- An optimal policy is one that maximizes expected utility if followed



## Values of States

- **Value function:**  $V^\pi(s)$  specifies value of following  $\pi$  starting in  $s$

$$V^\pi(s) = \mathbb{E} \left[ \sum_{t=1}^{\infty} \delta^{t-1} r_t(s_t, \pi(s_t), s_{t+1}) \mid s_0 = s \right]$$

- **State-action value function:**  $Q^\pi(s, a)$  returns value of starting in  $s$ , taking  $a$ , and then continuing according to  $\pi$
- These two are related to each other by

$$Q^\pi(s, a) = \sum_{s'} p(s, a, s') (r(s, a, s') + \delta V^\pi(s'))$$

$$V^\pi(s) = Q^\pi(s, \pi(s))$$

$$\Rightarrow V^\pi(s) = \sum_{s'} p(s, \pi(s), s') (r(s, \pi(s), s') + \delta V^\pi(s'))$$



# Policy Evaluation

---

---

Initialize  $V_0^\pi(s) \leftarrow 0$  for all states  $s$ ;

**for**  $t = 1 \dots T$  **do**

**for each state**  $s$  **do**

$V_t^\pi(s) \leftarrow \sum_{s'} p(s, \pi(s), s') (r(s, \pi(s), s') + \delta V_{t-1}^\pi(s'))$

---

- How many iterations should we have (what should  $T$  be)?
- Repeat until values do not change much:

$$\max_{s \in \mathcal{S}} |V_t^\pi(s) - V_{t-1}^\pi(s)| < \epsilon$$



## Solving MDP: Bellman Equations

$$Q^*(s, a) = \sum_{s'} p(s, a, s') (r(s, a, s') + \delta V^*(s'))$$

$$V^*(s) = \max_{a \in A} Q^*(s, a)$$

$$\Rightarrow V^*(s) = \max_{a \in A} \sum_{s'} p(s, a, s') (r(s, a, s') + \delta V^*(s'))$$



# Value Iteration

---

---

```
Initialize  $V_0(s) \leftarrow 0$  for all states  $s$ ;  
repeat until  $V(s)$  converges for all  $s$   
  for each state  $s$  do  
     $V_t(s) \leftarrow \max_{a \in A} \sum_{s'} p(s, a, s') (r(s, a, s') + \delta V_{t-1}(s'))$ 
```

---

- Bellman equations **characterize** the optimal values
- Value iteration **computes** them
- Value iteration is just a **fixed-point** solution method



# Policy Extraction

- Given  $V^*$ , we can compute optimal policy as follows:

$$\pi^*(s) = \operatorname{argmax}_a \sum_{s'} p(s, a, s') (r(s, a, s') + \delta V^*(s'))$$

- Given  $Q^*$ , we can compute optimal policy as follows:

$$\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$$

- Takeaway: actions are easier to select from Q-values than values!



# Problems with Value Iteration

- It is slow -  $O(S^2A)$
- The max at each state rarely changes
- The policy often converges long before the values





# Policy Iteration

- (I) **Policy evaluation:** Calculate values for some fixed policy
- (II) **Policy improvement:** Extract policy given these values
- Repeat steps until policy converges



# Value Iteration vs Policy Iteration

- Both compute the same thing (all optimal values)
- In **value iteration**:
  - Every iteration updates both values and (implicitly) policy
  - We don't track policy, but taking max over actions implicitly recomputes it
- In **policy iteration**:
  - We do several passes that update values with fixed policy (each pass is fast because we consider only one action, not all of them)
  - After policy is evaluated, a new policy is chosen (slow like a value iteration pass)
  - New policy will be better (or we're done)
- Both are **dynamic programs** for solving MDPs




# Outline

1. Markov Decision Processes
2. Definition of Stochastic Games
3. Strategies and Equilibria in Stochastic Games



# Stochastic Games (a.k.a. Markov Games): Introduction

- Lloyd Shapley introduced stochastic games in early 1950s
- Stochastic games generalize repeated games
  - Agents repeatedly play games from set of stage games
- Stochastic games generalize Markov decision process
  - Game at each step only depends on outcome of previous step
- **Single-state** stochastic game = (infinitely) repeated game
- **Single-agent** stochastic game = MDP




Lloyd Shapley (1923–2016)




# Repeated Games vs Stochastic Games

## Repeated Games



## Stochastic Games



# Stochastic Games: Formal Definition<sup>1</sup>

- $S$  is finite set of **stage games**
- $N$  is finite set of  $n$  agents
- $A_i$  is finite set of actions available to agent  $i$
- $p : S \times A \times S \mapsto [0, 1]$  is **transition probability function**
  - $p(s, a, s')$  is probability of going from  $s$  to  $s'$  after action profile  $a$
- $r_i : S \times A \mapsto \mathbb{R}$  is real-valued **utility function** for agent  $i$ 
  - $r_i(s, a)$  is agent  $i$ 's utility at state  $s$  for action profile  $a$

---

<sup>1</sup>Note that this definition assume actions available to agents are the same across different stage games. Changing this assumption leads to a more involved notation.



# Outline

1. Markov Decision Processes
2. Definition of Stochastic Games
3. Strategies and Equilibria in Stochastic Games



# Stochastic Games: Strategies

- Let  $h_t = (s_0, a_0, s_1, a_1, \dots, a_{t-1}, s_t)$  denote **history** of  $t$  stages
- Let  $H_t$  be set of all possible histories of this length
- Set of all **deterministic** strategies for agent  $i$  is

$$\prod_{t, H_t} A_i$$

- Agents' strategies can consist of any mixture over deterministic strategies
- However, there are several restricted classes of strategies





# Behavioral, Markov, and Stationary Strategies

- **Behavioral** strategy  $\pi_i(h_t, a_i)$  returns probability of playing  $a_i$  for  $h_t$ 
  - Mixing takes place at each history independently
- **Markov** strategy  $\pi_i$  is behavioral strategy s.t.  $\pi_i(h_t, a_i) = \pi_i(h'_t, a_i)$  if  $s_t = s'_t$ 
  - $s_t$  and  $s'_t$  are final states of  $h_t$  and  $h'_t$ , respectively
  - For each  $t$ , distribution over actions depends only on current state
- **Stationary** strategy  $\pi_i$  is a Markov strategy s.t.  $\pi_i(h_{t_1}, a_i) = \pi_i(h'_{t_2}, a_i)$  if  $s_{t_1} = s'_{t_2}$ 
  - $s_{t_1}$  and  $s'_{t_2}$  are final states of  $h_{t_1}$  and  $h'_{t_2}$ , respectively
  - This removes possible dependence on time  $t$



# Markov-perfect Equilibrium (MPE)

- Strategy  $\pi$  is **MPE** if it is Markov strategy and is NE regardless of starting state

$$V_i^\pi(s) \geq V_i^{(\pi'_i, \pi_{-i})}(s) \quad \forall i, s, \pi'_i$$

- MPE is similar to **subgame-perfect equilibrium** in perfect-information games
- Every  $n$ -player, general-sum, **discounted-reward** stochastic game has MPE



# Computing Equilibrium

- Poly-time algorithms are not generally available for full class of stochastic games
- However, they exist for several nontrivial sub-classes
- E.g., 2-player, general-sum, discounted-reward, **single-controller** stochastic games
  - Transitions depend on single agent: if  $a_i = a'_i$ , then  $p(s, a, s') = p(s, a', s') \forall s, s'$
- E.g., 2-player, general-sum, discounted-reward, **separable-reward, state-independent-transition (SR-SIT)** stochastic games
  - $r_i(s, a) = f(s) + g(a) \forall i, s, a$ , and
  - $p(s, a, s'') = p(s', a, s'') \forall s, s', s'', a$
- E.g., 2-player, zero-sum, discounted-reward stochastic games



## Shapley Algorithm: Finding MPE in 2-player Zero-sum Games

---

---

Initialize  $V_0(s)$  arbitrarily for all  $s$ ;

▷ Agent 1's utility for being in  $s$

**repeat** until  $V(s)$  converges for all  $s$

**for** each state  $s$  **do**

    Compute matrix game  $G(s, V_{t-1})$ :

$$u(s, a) = r(s, a) + \delta \sum_{s'} p(s, a, s') V_{t-1}(s')$$


**for** each state  $s$  **do**

$$V_t(s) \leftarrow \max_{\pi_1} \min_{\pi_2} u_1(s, \pi_1, \pi_2)$$

- 
- Shapley's algorithm is **extension of value iteration** to stochastic games



# Shapley Algorithm: Example



$$-3 + \delta (0.7 \times 2 + 0.3 \times 5) \\ = -3 + 2.9\delta$$

	$-3 + 2.9\delta$

$G(s_1, V_0)$



## Pollatschek & Avi-Itzhak Algorithm (Extension of Policy Iteration)

---

Initialize  $V(s)$  arbitrarily for all  $s$ ; ▷ Agent 1's utility for being in  $s$   
**repeat** until  $\pi_1(s)$  and  $\pi_2(s)$  converge for all  $s$   
  **for** each state  $s$  **do**  
    Compute matrix game  $G(s, V)$  as in Shapley's algorithm;  
     $\pi_1(s) \leftarrow$  maxmin strategy of Agent 1 in  $G(s, V)$ ;  
     $\pi_2(s) \leftarrow$  minmax strategy agents Agent 1 in  $G(s, V)$ ;  
  Calculate  $V(s)$  with policy evaluation for  $\pi_1$  and  $\pi_2$

---



# Acknowledgment

- This lecture is a slightly modified version of ones prepared by
  - Dan Klein and Pieter Abbeel [UC Berkeley CS 188]
  - Vincent Conitzer [Duke CPS 590.4]

