

Game-theoretic Foundations of Multi-agent Systems

Lecture 7: Stochastic Games

Seyed Majid Zahedi



Outline

1. Markov Decision Processes
2. Definition of Stochastic Games
3. Strategies and Equilibria in Stochastic Games




Motivation: Non-deterministic Search in Grid World




Grid World Actions

Deterministic




Grid World Actions

Deterministic




Stochastic



A Grid World Instance

- Agent lives in a grid
- Walls block agent's path
- Actions do not always go as planned
 - 80% of time, action "North" takes us north
 - 10% of time, "North" takes us west; 10% east
 - If there is a wall in the direction the agent would have been taken, the agent stays put
- Agent receives rewards at each step
- Goal is to maximize sum of rewards




Change in Notation

- So far, we have used s to denote **strategy profile**
- In this lecture, we use π for strategy
- We use s to denote **state**



Markov Property

- Given present state, future and past are independent




Andrey Markov (1856-1922)



Markov Property

- Given present state, future and past are independent
- Future state depends only on current state and action




Andrey Markov (1856-1922)



Markov Property

- Given present state, future and past are independent
- Future state depends only on current state and action



Andrey Markov (1856-1922)

$$P(S_{t+1} = s \mid S_t = s_t, A_t = a_t, S_{t-1} = s_{t-1}, A_{t-1} = a_{t-1}, \dots, S_0 = s_0) = \\ P(S_{t+1} = s \mid S_t = s_t, A_t = a_t)$$

Markov Decision Processes: Formal Definition

- S is set of states and A is set of actions



Markov Decision Processes: Formal Definition

- S is set of states and A is set of actions
- $p : S \times A \times S \mapsto [0, 1]$ specifies transition probabilities
 - $p(s, a, s')$ is probability of going to s' when taking a in s



Markov Decision Processes: Formal Definition

- S is set of states and A is set of actions
- $p : S \times A \times S \mapsto [0, 1]$ specifies transition probabilities
 - $p(s, a, s')$ is probability of going to s' when taking a in s
- $r : S \times A \times S \mapsto \mathbb{R}$ returns reward
 - $r(s, a, s')$ is reward of going to s' when taking a in s



Markov Decision Processes: Formal Definition

- S is set of states and A is set of actions
- $p : S \times A \times S \mapsto [0, 1]$ specifies transition probabilities
 - $p(s, a, s')$ is probability of going to s' when taking a in s
- $r : S \times A \times S \mapsto \mathbb{R}$ returns reward
 - $r(s, a, s')$ is reward of going to s' when taking a in s
- There are two ways to aggregate rewards
 - Limit-average reward: $\lim_{T \rightarrow \infty} \sum_{t=1}^T r_t / T$
 - Future-discounted reward: $\sum_{t=1}^{\infty} \delta^{t-1} r_t$



Markov Decision Processes: Formal Definition

- S is set of states and A is set of actions
- $p : S \times A \times S \mapsto [0, 1]$ specifies transition probabilities
 - $p(s, a, s')$ is probability of going to s' when taking a in s
- $r : S \times A \times S \mapsto \mathbb{R}$ returns reward
 - $r(s, a, s')$ is reward of going to s' when taking a in s
- There are two ways to aggregate rewards
 - Limit-average reward: $\lim_{T \rightarrow \infty} \sum_{t=1}^T r_t / T$
 - Future-discounted reward: $\sum_{t=1}^{\infty} \delta^{t-1} r_t$




Markov Decision Processes: Formal Definition

- S is set of states and A is set of actions
- $p : S \times A \times S \mapsto [0, 1]$ specifies transition probabilities
 - $p(s, a, s')$ is probability of going to s' when taking a in s
- $r : S \times A \times S \mapsto \mathbb{R}$ returns reward
 - $r(s, a, s')$ is reward of going to s' when taking a in s
- There are two ways to aggregate rewards
 - Limit-average reward: $\lim_{T \rightarrow \infty} \sum_{t=1}^T r_t / T$
 - **Future-discounted reward:** $\sum_{t=1}^{\infty} \delta^{t-1} r_t$
(we will focus on this)




Example: Racing

- Autonomous car wants to travel far, quickly
- There are 3 states: Cool, Warm, Overheated, and 2 actions: **Slow**, **Fast**
- Going faster gets double reward



Policies

- A (stationary, deterministic) policy $\pi : S \mapsto A$ gives action for each state
- An optimal policy is one that maximizes expected utility if followed



Values of States

- **Value function:** $V^\pi(s)$ specifies value of following π starting in s

$$V^\pi(s) = \mathbb{E} \left[\sum_{t=1}^{\infty} \delta^{t-1} r_t(s_t, \pi(s_t), s_{t+1}) \mid s_0 = s \right]$$



Values of States

- **Value function:** $V^\pi(s)$ specifies value of following π starting in s

$$V^\pi(s) = \mathbb{E} \left[\sum_{t=1}^{\infty} \delta^{t-1} r_t(s_t, \pi(s_t), s_{t+1}) \mid s_0 = s \right]$$

- **State-action value function:** $Q^\pi(s, a)$ returns value of starting in s , taking a , and then continuing according to π



Values of States

- **Value function:** $V^\pi(s)$ specifies value of following π starting in s

$$V^\pi(s) = \mathbb{E} \left[\sum_{t=1}^{\infty} \delta^{t-1} r_t(s_t, \pi(s_t), s_{t+1}) \mid s_0 = s \right]$$

- **State-action value function:** $Q^\pi(s, a)$ returns value of starting in s , taking a , and then continuing according to π
- These two are related to each other by

$$Q^\pi(s, a) = \sum_{s'} p(s, a, s') (r(s, a, s') + \delta V^\pi(s'))$$

$$V^\pi(s) = Q^\pi(s, \pi(s))$$



Values of States

- **Value function:** $V^\pi(s)$ specifies value of following π starting in s

$$V^\pi(s) = \mathbb{E} \left[\sum_{t=1}^{\infty} \delta^{t-1} r_t(s_t, \pi(s_t), s_{t+1}) \mid s_0 = s \right]$$

- **State-action value function:** $Q^\pi(s, a)$ returns value of starting in s , taking a , and then continuing according to π
- These two are related to each other by

$$Q^\pi(s, a) = \sum_{s'} p(s, a, s') (r(s, a, s') + \delta V^\pi(s'))$$

$$V^\pi(s) = Q^\pi(s, \pi(s))$$

$$\Rightarrow V^\pi(s) = \sum_{s'} p(s, \pi(s), s') (r(s, \pi(s), s') + \delta V^\pi(s'))$$



Policy Evaluation

Initialize $V_0^\pi(s) \leftarrow 0$ for all states s ;

for $t = 1 \dots T$ **do**

for each state s **do**

$V_t^\pi(s) \leftarrow \sum_{s'} p(s, \pi(s), s') (r(s, \pi(s), s') + \delta V_{t-1}^\pi(s'))$

- How many iterations should we have (what should T be)?



Policy Evaluation

Initialize $V_0^\pi(s) \leftarrow 0$ for all states s ;

for $t = 1 \dots T$ **do**

for each state s **do**

$$\quad \quad V_t^\pi(s) \leftarrow \sum_{s'} p(s, \pi(s), s') (r(s, \pi(s), s') + \delta V_{t-1}^\pi(s'))$$

- How many iterations should we have (what should T be)?
- Repeat until values do not change much:

$$\max_{s \in S} |V_t^\pi(s) - V_{t-1}^\pi(s)| < \epsilon$$

Solving MDP: Bellman Equations

$$Q^*(s, a) = \sum_{s'} p(s, a, s') (r(s, a, s') + \delta V^*(s'))$$

$$V^*(s) = \max_{a \in A} Q^*(s, a)$$

$$\Rightarrow V^*(s) = \max_{a \in A} \sum_{s'} p(s, a, s') (r(s, a, s') + \delta V^*(s'))$$



Value Iteration

Initialize $V_0(s) \leftarrow 0$ for all states s ;

repeat until $V(s)$ converges for all s

for each state s **do**

$V_t(s) \leftarrow \max_{a \in A} \sum_{s'} p(s, a, s') (r(s, a, s') + \delta V_{t-1}(s'))$



Value Iteration

Initialize $V_0(s) \leftarrow 0$ for all states s ;

repeat until $V(s)$ converges for all s

for each state s **do**

$$V_t(s) \leftarrow \max_{a \in A} \sum_{s'} p(s, a, s') (r(s, a, s') + \delta V_{t-1}(s'))$$

- Bellman equations **characterize** the optimal values



Value Iteration

Initialize $V_0(s) \leftarrow 0$ for all states s ;

repeat *until* $V(s)$ converges for all s

for *each* state s **do**

$$V_t(s) \leftarrow \max_{a \in A} \sum_{s'} p(s, a, s') (r(s, a, s') + \delta V_{t-1}(s'))$$

-
- Bellman equations **characterize** the optimal values
 - Value iteration **computes** them



Value Iteration

Initialize $V_0(s) \leftarrow 0$ for all states s ;

repeat *until* $V(s)$ converges for all s

for *each state* s **do**

$V_t(s) \leftarrow \max_{a \in A} \sum_{s'} p(s, a, s') (r(s, a, s') + \delta V_{t-1}(s'))$

- Bellman equations **characterize** the optimal values
- Value iteration **computes** them
- Value iteration is just a **fixed-point** solution method



Policy Extraction

- Given V^* , we can compute optimal policy as follows:

$$\pi^*(s) = \operatorname{argmax}_a \sum_{s'} p(s, a, s') (r(s, a, s') + \delta V^*(s'))$$



Policy Extraction

- Given V^* , we can compute optimal policy as follows:

$$\pi^*(s) = \operatorname{argmax}_a \sum_{s'} p(s, a, s') (r(s, a, s') + \delta V^*(s'))$$

- Given Q^* , we can compute optimal policy as follows:

$$\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$$



Policy Extraction

- Given V^* , we can compute optimal policy as follows:

$$\pi^*(s) = \operatorname{argmax}_a \sum_{s'} p(s, a, s') (r(s, a, s') + \delta V^*(s'))$$

- Given Q^* , we can compute optimal policy as follows:

$$\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$$

- Takeaway: actions are easier to select from Q-values than values!

Problems with Value Iteration

- It is slow - $O(S^2A)$
- The max at each state rarely changes
- The policy often converges long before the values



Policy Iteration

- (I) **Policy evaluation:** Calculate values for some fixed policy



Policy Iteration

- (I) **Policy evaluation:** Calculate values for some fixed policy
- (II) **Policy improvement:** Extract policy given these values



Policy Iteration

- (I) **Policy evaluation:** Calculate values for some fixed policy
- (II) **Policy improvement:** Extract policy given these values
- Repeat steps until policy converges



Value Iteration vs Policy Iteration

- Both compute the same thing (all optimal values)



Value Iteration vs Policy Iteration

- Both compute the same thing (all optimal values)
- In **value iteration**:



Value Iteration vs Policy Iteration

- Both compute the same thing (all optimal values)
- In **value iteration**:
 - Every iteration updates both values and (implicitly) policy



Value Iteration vs Policy Iteration

- Both compute the same thing (all optimal values)
- In **value iteration**:
 - Every iteration updates both values and (implicitly) policy
 - We don't track policy, but taking max over actions implicitly recomputes it



Value Iteration vs Policy Iteration

- Both compute the same thing (all optimal values)
- In **value iteration**:
 - Every iteration updates both values and (implicitly) policy
 - We don't track policy, but taking max over actions implicitly recomputes it
- In **policy iteration**:



Value Iteration vs Policy Iteration

- Both compute the same thing (all optimal values)
- In **value iteration**:
 - Every iteration updates both values and (implicitly) policy
 - We don't track policy, but taking max over actions implicitly recomputes it
- In **policy iteration**:
 - We do several passes that update values with fixed policy (each pass is fast because we consider only one action, not all of them)



Value Iteration vs Policy Iteration

- Both compute the same thing (all optimal values)
- In **value iteration**:
 - Every iteration updates both values and (implicitly) policy
 - We don't track policy, but taking max over actions implicitly recomputes it
- In **policy iteration**:
 - We do several passes that update values with fixed policy (each pass is fast because we consider only one action, not all of them)
 - After policy is evaluated, a new policy is chosen (slow like a value iteration pass)



Value Iteration vs Policy Iteration

- Both compute the same thing (all optimal values)
- In **value iteration**:
 - Every iteration updates both values and (implicitly) policy
 - We don't track policy, but taking max over actions implicitly recomputes it
- In **policy iteration**:
 - We do several passes that update values with fixed policy (each pass is fast because we consider only one action, not all of them)
 - After policy is evaluated, a new policy is chosen (slow like a value iteration pass)
 - New policy will be better (or we're done)



Value Iteration vs Policy Iteration

- Both compute the same thing (all optimal values)
- In **value iteration**:
 - Every iteration updates both values and (implicitly) policy
 - We don't track policy, but taking max over actions implicitly recomputes it
- In **policy iteration**:
 - We do several passes that update values with fixed policy (each pass is fast because we consider only one action, not all of them)
 - After policy is evaluated, a new policy is chosen (slow like a value iteration pass)
 - New policy will be better (or we're done)
- Both are **dynamic programs** for solving MDPs


Outline

1. Markov Decision Processes
2. Definition of Stochastic Games
3. Strategies and Equilibria in Stochastic Games



Stochastic Games (a.k.a. Markov Games): Introduction

- Lloyd Shapley introduced stochastic games in early 1950s




Lloyd Shapley (1923–2016)



Stochastic Games (a.k.a. Markov Games): Introduction

- Lloyd Shapley introduced stochastic games in early 1950s
- Stochastic games generalize repeated games




Lloyd Shapley (1923-2016)



Stochastic Games (a.k.a. Markov Games): Introduction

- Lloyd Shapley introduced stochastic games in early 1950s
- Stochastic games generalize repeated games
 - Agents repeatedly play games from set of stage games




Lloyd Shapley (1923-2016)



Stochastic Games (a.k.a. Markov Games): Introduction

- Lloyd Shapley introduced stochastic games in early 1950s
- Stochastic games generalize repeated games
 - Agents repeatedly play games from set of stage games
- Stochastic games generalize Markov decision process




Lloyd Shapley (1923-2016)



Stochastic Games (a.k.a. Markov Games): Introduction

- Lloyd Shapley introduced stochastic games in early 1950s
- Stochastic games generalize repeated games
 - Agents repeatedly play games from set of stage games
- Stochastic games generalize Markov decision process
 - Game at each step only depends on outcome of previous step




Lloyd Shapley (1923-2016)



Stochastic Games (a.k.a. Markov Games): Introduction

- Lloyd Shapley introduced stochastic games in early 1950s
- Stochastic games generalize repeated games
 - Agents repeatedly play games from set of stage games
- Stochastic games generalize Markov decision process
 - Game at each step only depends on outcome of previous step
- **Single-state** stochastic game = (infinitely) repeated game




Lloyd Shapley (1923–2016)



Stochastic Games (a.k.a. Markov Games): Introduction


- Lloyd Shapley introduced stochastic games in early 1950s
- Stochastic games generalize repeated games
 - Agents repeatedly play games from set of stage games
- Stochastic games generalize Markov decision process
 - Game at each step only depends on outcome of previous step
- **Single-state** stochastic game = (infinitely) repeated game
- **Single-agent** stochastic game = MDP



Lloyd Shapley (1923-2016)


Repeated Games vs Stochastic Games

Repeated Games




Repeated Games vs Stochastic Games

Repeated Games



Stochastic Games



Stochastic Games: Formal Definition¹

- S is finite set of **stage games**
- N is finite set of n agents
- A_i is finite set of actions available to agent i
- $p : S \times A \times S \mapsto [0, 1]$ is **transition probability function**
 - $p(s, a, s')$ is probability of going from s to s' after action profile a
- $r_i : S \times A \mapsto \mathbb{R}$ is real-valued **utility function** for agent i
 - $r_i(s, a)$ is agent i 's utility at state s for action profile a

¹Note that this definition assume actions available to agents are the same across different stage games. Changing this assumption leads to a more involved notation.



Outline

1. Markov Decision Processes
2. Definition of Stochastic Games
3. Strategies and Equilibria in Stochastic Games



Stochastic Games: Strategies

- Let $h_t = (s_0, a_0, s_1, a_1, \dots, a_{t-1}, s_t)$ denote **history** of t stages
- Let H_t be set of all possible histories of this length
- Set of all **deterministic** strategies for agent i is

$$\prod_{t, H_t} A_i$$

- Agents' strategies can consist of any mixture over deterministic strategies
- However, there are several restricted classes of strategies



Behavioral, Markov, and Stationary Strategies

- **Behavioral** strategy $\pi_i(h_t, a_i)$ returns probability of playing a_i for h_t
 - Mixing takes place at each history independently



Behavioral, Markov, and Stationary Strategies

- **Behavioral** strategy $\pi_i(h_t, a_i)$ returns probability of playing a_i for h_t
 - Mixing takes place at each history independently
- **Markov** strategy π_i is behavioral strategy s.t. $\pi_i(h_t, a_i) = \pi_i(h'_t, a_i)$ if $s_t = s'_t$
 - s_t and s'_t are final states of h_t and h'_t , respectively
 - For each t , distribution over actions depends only on current state



Behavioral, Markov, and Stationary Strategies

- **Behavioral** strategy $\pi_i(h_t, a_i)$ returns probability of playing a_i for h_t
 - Mixing takes place at each history independently
- **Markov** strategy π_i is behavioral strategy s.t. $\pi_i(h_t, a_i) = \pi_i(h'_t, a_i)$ if $s_t = s'_t$
 - s_t and s'_t are final states of h_t and h'_t , respectively
 - For each t , distribution over actions depends only on current state
- **Stationary** strategy π_i is a Markov strategy s.t. $\pi_i(h_{t_1}, a_i) = \pi_i(h'_{t_2}, a_i)$ if $s_{t_1} = s'_{t_2}$
 - s_{t_1} and s'_{t_2} are final states of h_{t_1} and h'_{t_2} , respectively
 - This removes possible dependence on time t

Markov-perfect Equilibrium (MPE)

- Strategy π is **MPE** if it is Markov strategy and is NE regardless of starting state

$$V_i^\pi(s) \geq V_i^{(\pi'_i, \pi_{-i})}(s) \quad \forall i, s, \pi'_i$$

- MPE is similar to **subgame-perfect equilibrium** in perfect-information games
- Every n -player, general-sum, **discounted-reward** stochastic game has MPE



Computing Equilibrium

- Poly-time algorithms are not generally available for full class of stochastic games



Computing Equilibrium

- Poly-time algorithms are not generally available for full class of stochastic games
- However, they exist for several nontrivial sub-classes



Computing Equilibrium

- Poly-time algorithms are not generally available for full class of stochastic games
- However, they exist for several nontrivial sub-classes
- E.g., 2-player, general-sum, discounted-reward, **single-controller** stochastic games
 - Transitions depend on single agent: if $a_i = a'_i$, then $p(s, a, s') = p(s, a', s') \forall s, s'$



Computing Equilibrium

- Poly-time algorithms are not generally available for full class of stochastic games
- However, they exist for several nontrivial sub-classes
- E.g., 2-player, general-sum, discounted-reward, **single-controller** stochastic games
 - Transitions depend on single agent: if $a_i = a'_i$, then $p(s, a, s') = p(s, a', s') \forall s, s'$
- E.g., 2-player, general-sum, discounted-reward, **separable-reward, state-independent-transition (SR-SIT)** stochastic games
 - $r_i(s, a) = f(s) + g(a) \forall i, s, a$, and
 - $p(s, a, s'') = p(s', a, s'') \forall s, s', s'', a$



Computing Equilibrium

- Poly-time algorithms are not generally available for full class of stochastic games
- However, they exist for several nontrivial sub-classes
- E.g., 2-player, general-sum, discounted-reward, **single-controller** stochastic games
 - Transitions depend on single agent: if $a_i = a'_i$, then $p(s, a, s') = p(s, a', s') \forall s, s'$
- E.g., 2-player, general-sum, discounted-reward, **separable-reward, state-independent-transition (SR-SIT)** stochastic games
 - $r_i(s, a) = f(s) + g(a) \forall i, s, a$, and
 - $p(s, a, s'') = p(s', a, s'') \forall s, s', s'', a$
- E.g., 2-player, zero-sum, discounted-reward stochastic games

Shapley Algorithm: Finding MPE in 2-player Zero-sum Games

Initialize $V_0(s)$ arbitrarily for all s ;

▷ Agent 1's utility for being in s

repeat until $V(s)$ converges for all s

for each state s **do**

 Compute matrix game $G(s, V_{t-1})$:

$$u(s, a) = r(s, a) + \delta \sum_{s'} p(s, a, s') V_{t-1}(s')$$

for each state s **do**

$$V_t(s) \leftarrow \max_{\pi_1} \min_{\pi_2} u_1(s, \pi_1, \pi_2)$$



Shapley Algorithm: Finding MPE in 2-player Zero-sum Games

Initialize $V_0(s)$ arbitrarily for all s ;

▷ Agent 1's utility for being in s

repeat until $V(s)$ converges for all s

for each state s **do**

 Compute matrix game $G(s, V_{t-1})$:

$$u(s, a) = r(s, a) + \delta \sum_{s'} p(s, a, s') V_{t-1}(s')$$

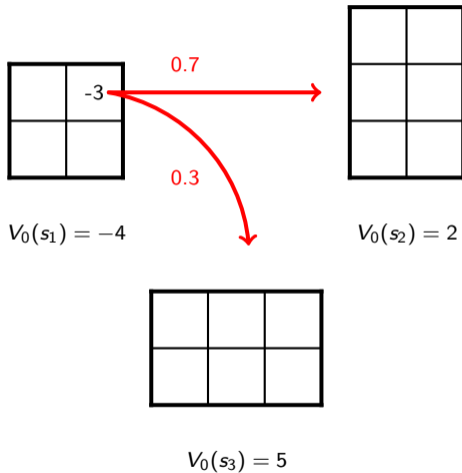
for each state s **do**

$$V_t(s) \leftarrow \max_{\pi_1} \min_{\pi_2} u_1(s, \pi_1, \pi_2)$$


-
- Shapley's algorithm is **extension of value iteration** to stochastic games



Shapley Algorithm: Example



Shapley Algorithm: Example



$$-3 + \delta (0.7 \times 2 + 0.3 \times 5) \\ = -3 + 2.9\delta$$

	$-3 + 2.9\delta$

$G(s_1, V_0)$



Pollatschek & Avi-Itzhak Algorithm (Extension of Policy Iteration)

Initialize $V(s)$ arbitrarily for all s ; ▷ Agent 1's utility for being in s
repeat until $\pi_1(s)$ and $\pi_2(s)$ converge for all s
 for each state s **do**
 Compute matrix game $G(s, V)$ as in Shapley's algorithm;
 $\pi_1(s) \leftarrow$ maxmin strategy of Agent 1 in $G(s, V)$;
 $\pi_2(s) \leftarrow$ minmax strategy agents Agent 1 in $G(s, V)$;
 Calculate $V(s)$ with policy evaluation for π_1 and π_2



Acknowledgment

- This lecture is a slightly modified version of ones prepared by
 - Dan Klein and Pieter Abbeel [UC Berkeley CS 188]
 - Vincent Conitzer [Duke CPS 590.4]

