

Game-theoretic Foundations of Multi-agent Systems

Lecture 9: Learning in Games

Seyed Majid Zahedi



Outline

1. Introduction
2. Fictitious Play
3. Best-response Dynamics
4. No-regret Learning
5. Background: Single-agent Reinforcement Learning
6. Multi-agent Reinforcement Learning



Single-agent vs Multi-agent Learning

- In artificial intelligence (AI), learning is usually performed by **single agent**
- Learning agent learns to function successfully in **unknown environment**
- In multi-agent setting, environment contains other agents
- Agents' learning changes the environment
- These changes **depend** in part on actions of learning agents
- Learning of each agent is **impacted** by learning performed by others
- Different learning rules lead to different **dynamical system**
- Simple learning rules can lead to complex global behaviors of system



Learning and Teaching

- In multi-agent systems, learning and teaching are inseparable
- Agents must consider what they have **learned** from others' past behavior
- They also must consider how they wish to **influence** others' future behavior
- In such setting, learning as **accumulating knowledge** is not always beneficial
- Accumulating knowledge should never hurt, one can always ignore what is learned
- But when one pre-commits to particular strategy for acting on accumulated knowledge, sometimes less is more
- E.g., in game of Chicken, if your opponent is learning your strategy to play best response, then optimal strategy is to always dare



Is Agent Learning in Optimal Way?

- In (repeated or stochastic) zero-sum games, this question is meaningful to ask
- In general, answer depends not only on learning procedure but also on others' behavior
- When all agents adopt same strategy, the setting is called **self-play**
 - E.g., all agent adopt TfT, or all adopt reinforcement learning (RL)
- One way to evaluate learning procedures is based on their performance in self-play
- But learning agents can also be judged by how they do in context of other agent types
 - TfT agent may perform well against TfT agents, but less well against RL agents
- Note that in GT, **optimal strategy** is replaced by **best response** (and equilibrium)



Properties of Learning Rules

- **Safety**: Guarantee agents at least their maxmin value
- **Rationality**: Settle on best response to opponent's strategy whenever opponent settles on stationary strategy
 - Opponent adopts same mixed strategy each time, regardless of the past
- **No regret**: Yield payoff that is no less than payoff agent could have obtained by playing any pure strategy against any set of opponents (details later!)



Outline

1. Introduction
- 2. Fictitious Play**
3. Best-response Dynamics
4. No-regret Learning
5. Background: Single-agent Reinforcement Learning
6. Multi-agent Reinforcement Learning



Fictitious Play: Introduction

- What are agents learning about?
- Arguably, most plausible answer is strategies of others
- **Fictitious** play (FP), one of earliest learning rules, takes this approach
- FP was first introduced by G. W. Brown in 1951¹
- Brown imagined that agents would “**simulate**” the game in their mind and update their future play based on this simulation; hence name fictitious play
- In its current use, FP is misnomer, since each play of the game actually occurs

¹Brown, G. W. “Iterative solution of games by fictitious play.” 1951



Fictitious Play

- Two agents repeatedly play stage game G
- $\eta_i^t(a_{-i})$ denotes number of times agent i has observed a_{-i} before time t
- η_i^1 represents **fictitious past** and cannot be zero for all a_{-i}
- Agents assume that their opponent is using **stationary mixed strategy**
- Agents **update their beliefs** about this strategy at each step according to:

$$\mu_i^t(a_{-i}) = \frac{\eta_i^t(a_{-i})}{\sum_{a'_{-i}} \eta_i^t(a'_{-i})}$$

- μ_i^t is **empirical distribution** of past actions and is treated as mixed strategy
- Agents best-respond to their beliefs about opponent' strategy

$$a_i^{t+1} = \operatorname{argmax}_{a_i} u_i(a_i, \mu_i^t)$$



Fictitious Play: Example

- Consider the following coordination game

	L	R
U	3,3	0,0
D	4,0	1,1

- Note that this game is dominant solvable with unique NE of (D, R)
- Suppose that $\eta_1^1 = (3, 0)$ and $\eta_2^1 = (1, 2.5)$
- FP proceeds as follows:

Round	1's η	2's η	1's action	2's action
1	(3, 0)	(1, 2.5)	D	L
2	(4, 0)	(1, 3.5)	D	R
3	(4, 1)	(1, 4.5)	D	R
4	(4, 2)	(1, 5.5)	D	R



Fictitious Play: Discussion

- In FP, agents **do not** need to know anything about their opponent's utilities
- FP is somewhat paradoxical as agents assume stationary strategy for their opponent, yet no agent plays stationary strategy except when FP converges
- Even though FP is **belief based** it is also **myopic**
- I.e., agents maximize current utility without considering their future ones
- Agents do not learn **true model** that generates empirical frequencies
- In other words, they do not learn how their opponent is actually playing the game



Convergence of Fictitious Play to Pure Strategies

- Let $\{a^t\}$ be sequence of action profiles generated by FP for G
- Sequence **converges** to a^* if there exists T s.t. $a^t = a^*$ for all $t \geq T$
- a^* is called **steady state** or **absorbing state** of FP
- (I) If sequence converges to a^* , then a^* is pure-strategy NE of G
- (II) If for some t , $a^t = a^*$, where a^* is strict NE of G , then $a^\tau = a^*$ for all $\tau > t$



Proof

- (I) is straightforward, for (II), let $a^t = a^*$, we want to show that $a^{t+1} = a^*$
- First, note that we can write μ as:

$$\mu_i^{t+1} = (1 - \alpha)\mu_i^t + \alpha a_{-i}^t = (1 - \alpha)\mu_i^t + \alpha a_{-i}^*$$

here, abusing notation, a_{-i}^t denotes degenerate probability distribution and:

$$\alpha = \frac{1}{\sum_{a'_{-i}} \eta_i^t(a'_{-i}) + 1}$$

- By linearity of expected utility, we have for all a_i :

$$u_i(a_i, \mu_i^{t+1}) = (1 - \alpha)u_i(a_i, \mu_i^t) + \alpha u_i(a_i, a_{-i}^*)$$

- Since a_i^* maximizes both terms, it follows that it is played at $t + 1$



Convergence of Fictitious Play to Mixed Strategies

- Of course, one cannot guarantee that fictitious play always converges to NE
- In FP, agents only play pure strategies and pure-strategy NE may not exist
- While FP sequence may not converge, its empirical distribution may
- Sequence $\{a^t\}$ converges to s^* in time-average sense if for all i and a_i :

$$\lim_{T \rightarrow \infty} \frac{\sum_{t=1}^T \mathbb{1}(a_i^t = a_i)}{T} = s_i^*(a_i)$$

$\mathbb{1}(\cdot)$ denotes the indicator function

- If FP sequence converges to s^* in the time-average sense, then s^* is NE



Proof

- Suppose $\{a^t\}$ converges to s^* in time-average sense, but s^* is not NE
- There is some i , a'_i , and a_i with $s_i^*(a_i) > 0$ s.t. $u_i(a'_i, s_{-i}^*) > u_i(a_i, s_{-i}^*)$
- Choose ϵ s.t. $\epsilon < (u_i(a'_i, s_{-i}^*) - u_i(a_i, s_{-i}^*)) / 2$
- Choose T s.t. for all $t \geq T$, $|\mu_i^t(a_{-i}) - s_{-i}^*(a_{-i})| < \epsilon / \max_{a'} u_i(a')$ for all a_{-i}
- This is possible because $\mu_i^t(a_{-i}) \rightarrow s_{-i}^*(a_{-i})$ by assumption



Proof (cont.)

- Then, for any $t \geq T$, we have:

$$\begin{aligned}u_i(a_i, \mu_i^t) &= \sum_{a_{-i}} u_i(a_i, a_{-i}) \mu_i^t(a_{-i}) \\ &\leq \sum_{a_{-i}} u_i(a_i, a_{-i}) s_{-i}^*(a_{-i}) + \epsilon \\ &\leq \sum_{a_{-i}} u_i(a'_i, a_{-i}) s_{-i}^*(a_{-i}) - \epsilon \\ &\leq \sum_{a_{-i}} u_i(a'_i, a_{-i}) \mu_i^t(a_{-i}) = u_i(a'_i, \mu_i^t)\end{aligned}$$

- So after sufficiently large t , a_i is never played
- This implies that as $t \rightarrow \infty$, $\mu_i^t(a_i) \rightarrow 0$, which contradicts with $s_i^*(a_i) > 0$



Example: Matching Pennies

- Consider the matching-pennies game

	H	T
H	1, -1	-1, 1
T	-1, 1	1, -1

Round	1's η	2's η	1's action	2's action
1	(1.5, 2)	(2, 1.5)	T	T
2	(1.5, 3)	(2, 2.5)	T	H
3	(2.5, 3)	(2, 3.5)	T	H
4	(3.5, 3)	(2, 4.5)	H	H
5	(4.5, 3)	(3, 4.5)	H	H
6	(5.5, 3)	(4, 4.5)	H	H
7	(6.5, 3)	(5, 4.5)	H	T

- FP continues as deterministic cycle, time average converges to unique NE



Example: (Anti-)Coordination Game

- Note that if empirical distribution of actions converges to NE, there is no guarantee on distribution of played outcomes
- Consider the following coordination game

	A	B
A	1, 1	0, 0
B	0, 0	1, 1

- Note that this game is unique NE of $((0.5, 0.5), (0.5, 0.5))$

Round	1's η	2's η	1's action	2's action
1	(0.5, 0)	(0, 0.5)	A	B
2	(0.5, 1)	(1, 0.5)	B	A
3	(1.5, 1)	(1, 1.5)	A	B
4	(1.5, 2)	(2, 1.5)	B	A



General Fictitious Play Convergence

- Fictitious play converges in time-average sense for game G if:
 - G is zero-sum game
 - G is two-player game where each agent has at most two actions (2x2 games)
 - G is solvable by iterated strict dominance
 - G is **identical-interest game**, i.e., all agents have same payoff function
 - G is **potential game** (more on this later!)



Non-convergence of Fictitious Play

- Convergence of fictitious play **can not** be guaranteed in general
- Shapley showed that in modified rock-scissors-paper game, FP does not converge

	Rock	Paper	Scissors
Rock	0, 0	0, 1	1, 0
Paper	1, 0	0, 0	0, 1
Scissors	0, 1	1, 0	0, 0

- This game has unique NE: each agent mixes uniformly
- Suppose $\eta_1^1 = (1, 0, 0)$ and $\eta_2^1 = (0, 1, 0)$
- Shapley showed that play cycles among 6 (off-diagonal) profiles with periods of ever-increasing length, thus non-convergence



Smooth Fictitious Play (SFP)

- Instead of best-responding to beliefs, agents respond randomly, but somewhat proportional to their expected utility

$$s_i^t(a_i | \mu_i^t) = \frac{\exp(u_i(a_i, \mu_i^t)/\gamma)}{\sum_{a'_i} \exp(u_i(a'_i, \mu_i^t)/\gamma)}$$

- γ is called the **smoothing** parameter
- This is called **soft-max** policy
- Soft-max policy respects best replies, but leaves room for **exploration**
- If all agents use SFP with sufficiently small γ_i , empirical play converges to ϵ -CCE



Outline

1. Introduction
2. Fictitious Play
- 3. Best-response Dynamics**
4. No-regret Learning
5. Background: Single-agent Reinforcement Learning
6. Multi-agent Reinforcement Learning



Best-response Dynamics (BRD): Introduction

- Agents start playing **arbitrary** actions
- In arbitrary order, agents take turns updating their action
- Agent update their action only if doing so can improve their utility
- This is repeated until no agents wants to update their action

Initialize $a = (a_1, \dots, a_n)$ to be arbitrary action profile;
while *there exists i such that $a_i \notin \operatorname{argmax}_{a \in A_i} u_i(a, a_{-i})$* **do**
 Let a'_i be such that $u_i(a'_i, a_{-i}) > u_i(a, a_{-i})$;
 Set $a_i \leftarrow a'_i$;
return a



Best-response Dynamics: Discussion

- If BRD halts, it returns pure strategy Nash equilibrium
 - Every agent must be playing best response
- Does BRD always halt?
 - No: Consider matching pennies/Rock Paper Scissors



Example: Congestion Games

- N is set of n agents
- M is set of m resources
- A_i is set of actions available to agent i
 - a_i represents subset of resources that agent i chooses (i.e., $a_i \subseteq M$)
- ℓ_j is congestion cost function for resources $j \in M$
 - $\ell_j(k)$ represents cost of congestion on resource j when k agents choose j
- $n_j(a)$ is number of agents who choose resource j (i.e., $n_j(a) = |\{i \mid j \in a_i\}|$)
- $c_i(a) = \sum_{j \in a_i} \ell_j(n_j(a))$ is total cost of agent
- Agents minimize their total cost (instead of maximizing their total utility)



BRD in Congestion Games

- Consider **potential function** $\phi : A \rightarrow \mathbb{R}$:

$$\phi(a) = \sum_{j=1}^m \sum_{k=1}^{n_j(a)} \ell_j(k)$$

(Note: **not** social welfare)

- How does ϕ change in one round of BRD? Say i switches from a_i to $b_i \in A_i$
- Well... We know it must have decreased agent i 's cost:

$$\begin{aligned} \Delta c_i &\equiv c_i(b_i, a_{-i}) - c_i(a_i, a_{-i}) \\ &= \sum_{j \in b_i \setminus a_i} \ell_j(n_j(a) + 1) - \sum_{j \in a_i \setminus b_i} \ell_j(n_j(a)) < 0 \end{aligned}$$



BRD in Congestion Games (cont.)

$$\phi(a) = \sum_{j=1}^m \sum_{k=1}^{n_j(a)} \ell_j(k)$$

- Change in potential is:

$$\begin{aligned} \Delta\phi &\equiv \phi(b_i, a_{-i}) - \phi(a_i, a_{-i}) \\ &= \sum_{j \in b_i \setminus a_i} \ell_j(n_j(a) + 1) - \sum_{j \in a_i \setminus b_i} \ell_j(n_j(a)) \\ &= \Delta c_j \end{aligned}$$

- Since ϕ can take on only finitely many values, this cannot go on forever
- And hence BRD halts in congestion games ...
- Which proves the **existence** of pure strategy Nash equilibria!



Example: Load Balancing Games on Identical Servers

- n clients $i \in N$ schedule jobs of size $w_i > 0$ on m identical servers M
- Action space $A_i = M$ for each client
- For each server $j \in M$, load $\ell_j(a) = \sum_{i:a_i=j} w_i$
- Cost of client i is load of server that i chooses : $c_i(a) = \ell_{a_i}(a)$



Load Balancing Games on Identical Servers: Discussion

- *Almost* congestion game — but server costs depend on **which** clients choose them
- BRD converges in load balancing games on identical servers
- Load balancing games on identical servers have pure strategy NE



BRD in Load Balancing Games on Identical Servers

- Consider **potential function** ϕ as:

$$\phi(\mathbf{a}) = \frac{1}{2} \sum_{j=1}^m \ell_j(\mathbf{a})^2$$

- Suppose i switches from server j to server j' :

$$\begin{aligned} \Delta c_i(\mathbf{a}) &\equiv c_i(j', \mathbf{a}_{-i}) - c_i(j, \mathbf{a}_{-i}) \\ &= \ell_{j'}(\mathbf{a}) + w_i - \ell_j(\mathbf{a}) \\ &< 0 \end{aligned}$$



BRD in Load Balancing Games on Identical Servers (cont.)

$$\begin{aligned}\Delta\phi(\mathbf{a}) &\equiv \phi(j', \mathbf{a}_{-i}) - \phi(j, \mathbf{a}_{-i}) \\ &= \frac{1}{2} ((l_{j'}(\mathbf{a}) + w_i)^2 + (l_j(\mathbf{a}) - w_i)^2 - l_{j'}(\mathbf{a})^2 - l_j(\mathbf{a})^2) \\ &= \frac{1}{2} (2w_i l_{j'}(\mathbf{a}) + w_i^2 - 2w_i l_j(\mathbf{a}) + w_i^2) \\ &= w_i (l_{j'}(\mathbf{a}) + w_i - l_j(\mathbf{a})) \\ &= w_i \cdot \Delta c_i(\mathbf{a}) \\ &< 0\end{aligned}$$

Note: $\Delta c_i \neq \Delta\phi$



Potential Games

- $\phi : A \rightarrow \mathbb{R}_{\geq 0}$ is **exact potential function** for game G if for all $a, i, a_i,$ and b_i :

$$\phi(b_i, a_{-i}) - \phi(a_i, a_{-i}) = c_i(b_i, a_{-i}) - c_i(a_i, a_{-i})$$

- $\phi : A \rightarrow \mathbb{R}_{\geq 0}$ is **ordinal potential function** for game G if for all $a, i, a_i,$ and b_i :

$$(c_i(b_i, a_{-i}) - c_i(a_i, a_{-i}) < 0) \Rightarrow (\phi(b_i, a_{-i}) - \phi(a_i, a_{-i}) < 0)$$

(i.e. the change in utility is always equal **in sign** to the change in potential)

- BRD is **guaranteed** to converge in game G **iff** G has ordinal potential function



BRD and Potential Games

- We've already seen ordinal potential function \Rightarrow BRD converges
- Lets prove other direction
- Consider graph $G = (V, E)$
- Let each $a \in A$ be a vertex in G (i.e., $V = A$)
- Add directed edge (a, b) if it is possible to go from b to a by best-response move
 - I.e., if there is i such that $b = (b_i, a_{-i})$, and $c_i(b_i, a_{-i}) < c_i(a)$
- BRD can be viewed as traversing this graph
 - Start at arbitrary vertex a , and then traverse arbitrary outgoing edges



BRD and Potential Games (cont.)

- Nash Equilibria are the sinks in this graph
- Suppose BRD converges \Rightarrow there are no cycles in this graph
- So, from every vertex a there is some sink s that is reachable (why?)
- We construct potential function $\phi(a)$ for each vertex a
- $\phi(a)$ is length of **longest** finite path from a to any sink s
- We need: for any edge $a \rightarrow b$, $\phi(b) < \phi(a)$.
- Its true! $\phi(a) \geq \phi(b) + 1$. (why?)



Outline

1. Introduction
2. Fictitious Play
3. Best-response Dynamics
- 4. No-regret Learning**
5. Background: Single-agent Reinforcement Learning
6. Multi-agent Reinforcement Learning



Sequential Prediction: Stock-prediction Example

- Every day GME goes **up** or **down**
- Goal is to predict direction each day before market opens (to buy or short)
- Market can behave arbitrarily/adversarially
- So there is no way we can promise to do **well**
- However, we get **advice**



Expert Advice

- There are N **experts** who make predictions in T rounds
- At each round t , each expert i makes prediction $p_i^t \in \{U, D\}$
- Expertise is self proclaimed — no promise experts know what they're talking about
- We (algorithm) want to aggregate predictions, to make our own prediction p_A^t
- We learn true outcome o^t at the end of each round
- If we predicted incorrectly (i.e. $p_A^t \neq o^t$), then we **made a mistake**



Expert Advice (cont.)

- Goal is to after a while do (almost) as well as **best** expert in hindsight
- To make things easy, we assume for now that there is one **perfect** expert
- Perfect expert never makes mistakes (but we don't know who the expert is)
- Can we find strategy that is guaranteed to make at most $\log(N)$ mistakes?



The Halving Algorithm

Let $S^1 \leftarrow \{1, \dots, N\}$ be set of all experts;

for $t = 1$ *to* T **do**

 Predict with majority vote;

 Observe the true outcome o^t ;

 Eliminate all experts that made a mistake: $S^{t+1} = \{j \in S^t \mid p_j^t = o^t\}$;



The Halving Algorithm: Analysis

- Algorithm predicts with majority vote
- Every time it makes a mistake, at least half of remaining experts are eliminated
- Hence $|S^{t+1}| \leq |S^t|/2$
- On the other hand, perfect expert is never eliminated
- Hence $|S^t| \geq 1$ for all t
- Since $|S^1| = N$, this means there can be at most $\log N$ mistakes
- But what if no expert is perfect? Say the best expert makes OPT mistakes
- Can we find a way to make not too many more than OPT mistakes?



The Iterated Halving Algorithm

Let $S^1 \leftarrow \{1, \dots, N\}$ be the set of all experts;

for $t = 1$ *to* T **do**

if $|S^t| = 0$ **then**

 Reset: Set $S^t \leftarrow \{1, \dots, N\}$

 Predict with majority vote;

 Eliminate all experts that made a mistake: $S^{t+1} = \{i \in S^t \mid p_i^t = o^t\}$;



The Iterated Halving Algorithm: Analysis

- Whenever algorithm makes mistake, we eliminate half of experts
- So algorithm can make at most $\log N$ mistakes between any two resets
- But if we reset, it is because since last reset, **every** expert has made mistake
- In particular, between any two resets, **best** expert has made at least 1 mistake
- Algorithm makes at most $\log(N)(\text{OPT} + 1)$ mistakes
- Algorithm is wasteful in that every time we reset, we forget what we have learned!
- How about just **downweight** experts who make mistakes?



The Weighted Majority Algorithm

Set weights $w_i^1 \leftarrow 1$ for all experts i ;

for $t = 1$ *to* T **do**

 Predict with weighted majority vote;

 Down-weight experts who made mistakes: (i.e., if $p_i^t \neq o^t$, set $w_i^{t+1} \leftarrow w_i^t/2$)



The Weighted Majority Algorithm: Analysis

- Let M be total number of mistakes that algorithm makes
- Let $W^t = \sum_i w_i^t$ be total weight at step t
- When algorithm makes mistake, at least half of total weight is cut in half
- So: $W^{t+1} \leq (3/4)W^t$
- If algorithm makes M mistakes, $W^T \leq N \cdot (3/4)^M$
- Let i^* be the best expert, $W^T > w_{i^*}^T = (1/2)^{\text{OPT}}$, which gives:

$$(1/2)^{\text{OPT}} \leq W \leq N(3/4)^M \Rightarrow (4/3)^M \leq N \cdot 2^{\text{OPT}} \Rightarrow M \leq 2.4(\text{OPT} + \log(N))$$

- Algorithm makes at most $2.4(\text{OPT} + \log(N))$ mistakes
- $\log(N)$ is constant, so ratio of mistakes to OPT is 2.4 in limit – not great, but not bad



What Do We Want in an Algorithm?

- Make only $1\times$ as many mistakes as OPT in limit, rather than $2.4\times$
- Handle N distinct actions (separate action for each expert), not just up and down
- Handle arbitrary costs in $[0, 1]$ per expert per round, not just right and wrong



New Model/Algorithm

- In rounds $1, \dots, T$, algorithm **chooses some expert** i^t
- Each expert i **experiences loss**: $\ell_i^t \in [0, 1]$
- Algorithm experiences the loss of the expert it chooses: $\ell_A^t = \ell_{i^t}^t$
- Total loss of expert i is $L_i^T = \sum_{t=1}^T \ell_i^t$
- Total loss of algorithm is $L_A^T = \sum_{t=1}^T \ell_A^t$
- Goal is to obtain loss **“not much worse”** than that of the best expert: $\min_i L_i^T$



Multiplicative Weights (MW) Algorithm (a.w.a. Hedge Algorithm)

Set weights $w_i^1 \leftarrow 1$ for all experts i ;

for $t = 1$ to T **do**

 Let $W^t = \sum_{i=1}^N w_i^t$;

 Choose expert i with probability w_i^t / W^t ;

 For each i , set $w_i^{t+1} \leftarrow w_i^t \cdot \exp(-\epsilon \ell_i^t)$;

- Can be viewed as “smoothed” version of weighted majority algorithm
- Has parameter ϵ which controls how quickly it down-weights experts
- Is *randomized* — chooses experts w.p. proportional to their weights
- Can be used with alternative update: $w_i^{t+1} \leftarrow w_i^t \cdot (1 - \epsilon \ell_i^t)$



Multiplicative Weights Algorithm: Discussion

- For any sequence of losses, and any expert k :

$$\frac{1}{T} \mathbb{E}[L_{MW}^T] \leq \frac{1}{T} L_k^T + \epsilon + \frac{\ln(N)}{\epsilon \cdot T}$$

- In particular, setting $\epsilon = \sqrt{\ln(N)/T}$:

$$\frac{1}{T} \mathbb{E}[L_{MW}^T] \leq \frac{1}{T} \min_k L_k^T + 2\sqrt{\frac{\ln(N)}{T}}$$

- Average loss quickly approaches that of best expert **exactly**, at rate of $1/\sqrt{T}$
- This works for **arbitrary** sequence of losses (e.g., chosen adaptively by adversary)
- So we could use it to play games (**experts** \leftrightarrow **actions** and **losses** \leftrightarrow **costs**)



Recall: Minimax Theorem (John von Neumann, 1928)

In any finite, two-player, zero-sum game, in any NE, each agent receives a payoff that is equal to both their maxmin value and their minmax value

$$\max_{s_i} \min_{s_{-i}} u_i(s_i, s_{-i}) = \min_{s_{-i}} \max_{s_i} u_i(s_i, s_{-i})$$



Simple Proof for Minimax Theorem

- Scale utilities such that u_1 is in $[0, 1]$
- Write $v_1 = \min_{s_2} \max_{s_1} u_1(s_1, s_2)$ and $v_2 = \max_{s_1} \min_{s_2} u_1(s_1, s_2)$
- Suppose theorem were false: $v_1 = v_2 + \epsilon$ for some constant $\epsilon > 0$
- Suppose A1 and A2 repeatedly play against each other as follows
 - A2 uses MW algorithm: at round t , $s_2^t(a_2) = w_{a_2}^t / W^t$
 - A1 plays best response to A2's strategy: $s_1^t = \operatorname{argmax}_{s_1} u_1(s_1, s_2^t)$



Simple Proof for Minimax Theorem (cont.)

- For A2's MW algorithm, we have:

$$\frac{1}{T} \sum_{t=1}^T \mathbb{E}[u_1(a_1^t, a_2^t)] \leq \frac{1}{T} \min_{a_2} \sum_{t=1}^T u_1(a_1^t, a_2) + 2\sqrt{\frac{\log n}{T}}$$

- Let \bar{s}_1 be mixed strategy that puts weight $1/T$ on each action a_1^t , we have:

$$\frac{1}{T} \min_{a_2} \sum_{t=1}^T u_1(a_1^t, a_2) = \min_{a_2} \sum_{t=1}^T \frac{1}{T} u_1(a_1^t, a_2) = \min_{a_2} u_1(\bar{s}_1, a_2)$$

- By definition, we have: $\min_{a_2} u_1(\bar{s}_1, a_2) \leq \max_{s_1} \min_{a_2} u_1(s_1, a_2) = v_2$, and so:

$$\frac{1}{T} \sum_{t=1}^T \mathbb{E}[u_1(a_1^t, a_2^t)] \leq v_2 + 2\sqrt{\frac{\log n}{T}}$$



Simple Proof for Minimax Theorem (cont.)

- On the other hand, A1 best responds to A2's mixed strategy:

$$\begin{aligned}\frac{1}{T} \sum_{t=1}^T \mathbb{E}[u_1(a_1^t, a_2^t)] &= \frac{1}{T} \sum_{t=1}^T \max_{a_1} u_1(a_1, s_2^t) \\ &\geq \frac{1}{T} \sum_{t=1}^T \min_{s_2} \max_{a_1} u_1(a_1, s_2) = \frac{1}{T} \sum_{t=1}^T v_1 = v_1\end{aligned}$$

- Combining these inequalities, we get: $v_1 \leq v_2 + 2\sqrt{\log n/T}$
- Since $v_1 = v_2 + \epsilon$, we have: $\epsilon \leq 2\sqrt{\log n/T}$
- Taking T large enough leads to contradiction



External Regret

- Sequence a^1, \dots, a^T has **external regret** of $\Delta(T)$ if for every agent i and action a'_i :

$$\frac{1}{T} \sum_{t=1}^T u_i(a^t) \geq \frac{1}{T} \sum_{t=1}^T u_i(a'_i, a_{-i}) - \Delta(T)$$

- If $\Delta(T) = o_T(1)$, we say that sequence of action profiles has *no* external regret
- External regret measures regret to the best **fixed** action in hindsight
- If a^1, \dots, a^T has ϵ external regret, then distribution π that puts weight $1/T$ on each a^t (i.e., empirical distribution of actions) forms **ϵ -approximate CCE**

$$\mathbb{E}_{a \sim \pi}[u_i(a)] = \frac{1}{T} \sum_{t=1}^T u_i(a^t) \geq \frac{1}{T} \sum_{t=1}^T u_i(a'_i, a_{-i}) - \epsilon = \mathbb{E}_{a \sim \pi}[u_i(a'_i, a_{-i})] - \epsilon$$



No-(external-)regret Dynamics

- Suppose that all agents use MW algorithm to choose between k actions
- After T steps, sequence of outcomes has external regret of $\Delta(T) = 2\sqrt{\log k/T}$
- Empirical distribution of outcomes forms $\Delta(T)$ -approximate CCE
- For $T = 4 \log(k)/\epsilon^2$, distribution of outcomes converges to ϵ -approximate CCE



Swap Regret

- Sequence a^1, \dots, a^T has **swap regret** of $\Delta(T)$ if for every agent i and every **switching function** $F_i : A_i \rightarrow A_i$:

$$\frac{1}{T} \sum_{t=1}^T u_i(a^t) \geq \frac{1}{T} \sum_{t=1}^T u_i(F_i(a_i), a_{-i}) - \Delta(T)$$

- If $\Delta(T) = o_T(1)$, we say that sequence of action profiles has no swap regret
- This measures regret to counterfactual case where every action of particular type is swapped with different action in hindsight, separately for each action
- E.g., “Every time i bought Microsoft, i should have bought Apple, and every time i bought Google, i should have bought Comcast.”
- If a^1, \dots, a^T has ϵ swap regret, then distribution π that picks among a^1, \dots, a^T uniformly at random is ϵ -approximate correlated equilibrium



Generalization

- For any agent i , F_i , and $a \in A$, define **regret** as:

$$\text{Regret}_i(a, F_i) = u_i(F_i(a_i), a_{-i}) - u_i(a)$$

- F_i is **constant** switching function if $F_i(a_i) = F_i(a'_i)$ for all $a_i, a'_i \in A_i$
- π is CCE if for every agent i and every constant switching function F_i :

$$\mathbb{E}_{a \sim \pi}[\text{Regret}_i(a, F_i)] \leq 0$$

- π is CE if for every agent i and every switching function F_i :

$$\mathbb{E}_{a \sim \pi}[\text{Regret}_i(a, F_i)] \leq 0$$



How to Achieve No Swap Regret

- Define set of time steps that expert j is selected:

$$S_j = \{t : a_t = j\}$$

- Observation: To achieve no swap regret it would be sufficient that for every j :

$$\frac{1}{|S_j|} \sum_{t \in S_j} \ell_{a_t}^t \leq \frac{1}{|S_j|} \min_i \sum_{t \in S_j} \ell_i^t + \Delta(T)$$

- No **swap** regret = no **external** regret separately on each sequence of actions S_j
- Best switching function in hindsight = swapping each action j for best fixed action in hindsight over S_j
- Idea: Run k copies of PW, one responsible for each S_j




Algorithm Sketch for No Swap Regret

- Initialize k copies of MW algorithm one for each of k actions
- Let $q(i)_1^t, \dots, q(i)_k^t$ be distribution over experts for copy i at time t
- Combine these into single distribution over experts: p_1^t, \dots, p_k^t (details later!)
- Let $\ell_1^t, \dots, \ell_k^t$ be losses for experts at time t
- For copy i of MW algorithm, we **report** losses $p_i^t \ell_1^t, \dots, p_i^t \ell_k^t$
- I.e., to copy i , we report the true losses scaled by p_i^t



No-swap-regret Algorithm



No-swap-regret Algorithm: Analysis

- Expected cost of the master algorithm:

$$\frac{1}{T} \sum_{t=1}^T \sum_{i=1}^k p_i^t \cdot \ell_i^t \quad (1)$$

- Expected cost under switching function F

$$\frac{1}{T} \sum_{t=1}^T \sum_{i=1}^k p_i^t \cdot \ell_{F(i)}^t \quad (2)$$

- Goal: prove that (1) is at most (2) plus $\Delta(T) = o_T(1)$



No-swap-regret Algorithm: Analysis (cont.)

- Expected cost of M_j :

$$\frac{1}{T} \sum_{t=1}^T \sum_{i=1}^k q(j)_i^t (p_j^t \cdot \ell_i^t) \quad (3)$$

- M_j is no-regret algorithm, so its cost is at most:

$$\frac{1}{T} \sum_{t=1}^T p_j^t \cdot \ell_{F(j)}^t + \Delta(T) \quad (4)$$

for any any arbitrary F



No-swap-regret Algorithm: Analysis (cont.)

- Summing inequality between (3) and (4) over all copies:

$$\frac{1}{T} \sum_{t=1}^T \sum_{i=1}^k \sum_{j=1}^k q(j)_i^t (p_j^t \cdot \ell_i^t) \leq \frac{1}{T} \sum_{t=1}^T \sum_{j=1}^k p_j^t \cdot \ell_{F(j)}^t + k \cdot \Delta(T) \quad (5)$$

- Right-hand side is equal to (2)
- For left-hand side to be equal to (1), we need:

$$p_i^t = \sum_{j=1}^k p_j^t \cdot q(j)_i^t$$



Combining Distributions

$$p_i^t = \sum_{j=1}^k p_j^t \cdot q(j)_i^t$$

- These might be familiar as those defining **stationary distribution** of **Markov chain**
 - There are k states, probability of going to state i from j is $q(j)_i^t$
 - Stationary distribution over states is $(p_1^t \dots p_k^t)$
- These equations **always** have solution as probability distribution
- Crucial property: two ways of viewing the distribution over experts:
 - Each expert i is chosen with probability p_i^t or
 - W.p. p_j^t we select copy j and then select expert i w.p. $q(j)_i^t$



Regret Matching

- α^t : Average per-step reward received by agent up until time t
- $\alpha^t(a)$: Average per-period reward that would have been received up until time t had pure strategy a was played by agent, assuming others played the same
- **Regret** at time t for not having played a : $R^t(a) = \alpha^t(a) - \alpha^t$
- **Regret matching**: At time t , choose action a w.p. proportional to its regret:

$$s^t(a) = \frac{R^t(a)^+}{\sum_{a'} R^t(a')^+}$$




Outline

1. Introduction
2. Fictitious Play
3. Best-response Dynamics
4. No-regret Learning
- 5. Background: Single-agent Reinforcement Learning**
6. Multi-agent Reinforcement Learning




Reinforcement Learning

- Still assume MDP
 - Set of states $s \in S$
 - Set of actions $a \in A$
 - Model $p(s, a, s')$
 - Reward $r(s, a, s')$
- Still looking for policy $\pi(s)$
- New twist: we do not know p or r
- I.e. we do not know which states are good or what actions do
- Must actually try actions and states out to learn



Offline (MDPs) vs. Online (RL)



Offline solution



Online solution

Why Not Use Policy Evaluation?

- Simplified **Bellman updates** calculate V and Q for a fixed policy

$$V_t^\pi(s) \leftarrow \sum_{s'} p(s, \pi(s), s') (r(s, \pi(s), s') + \delta V_{t-1}^\pi(s'))$$

- This approach fully exploited connections between the states
- Unfortunately, we need p and r to do it!



Temporal Difference (TD) Learning

- Main idea: learn from every experience!
 - Update $V(s)$ each time we experience a transition (s, a, s', r)
 - Likely outcomes s' will contribute updates more often
- Temporal difference learning of values
 - Policy still fixed, still doing evaluation!
 - Move values toward value of whatever successor occurs: running average

$$\text{Sample of } V(s) : r(s, a, s') + \delta V^\pi(s')$$

$$\text{Update of } V(s) : V^\pi(s) \leftarrow (1 - \alpha)V^\pi(s) + \alpha (r(s, a, s') + \delta V^\pi(s'))$$

$$\text{Same update : } V^\pi(s) \leftarrow V^\pi(s) + \alpha (r(s, a, s') + \delta V^\pi(s') - V^\pi(s))$$



Problems with TD Value Learning


- TD value learning is model-free way to do policy evaluation
- It mimics Bellman updates with running sample averages
- However, if we want to turn values into (new) policy, we need p and r !

$$\begin{aligned}\pi(s) &= \operatorname{argmax}_a Q(s, a) \\ Q^\pi(s, a) &= \sum_{s'} p(s, a, s') (r(s, a, s') + \delta V(s'))\end{aligned}$$

- To solve this, we can learn Q-values instead of values
- This makes action selection model-free too!



Active Reinforcement Learning



Q-learning

- Q-Learning is sample-based Q-value iteration

$$Q_t(s, a) \leftarrow \sum_{s'} p(s, a, s') \left(r(s, a, s') + \delta \max_{a' \in A} Q_{t-1}(s', a') \right)$$

- We learn $Q(s, a)$ values as we go

$$\text{Sample : } r(s, a, s') + \delta \max_{a' \in A} Q(s', a')$$

$$\text{Update : } Q(s, a) \leftarrow (1 - \alpha_t)Q(s, a) + \alpha_t \left(r(s, a, s') + \delta \max_{a' \in A} Q(s', a') \right)$$



Q-learning Algorithm

repeat *until convergence*

observe current state s ;

select action a and take it (e.g., via ϵ -greedy policy);

observe next state s' and reward $r(s, a, s')$;

$Q_{t+1}(s, a) \leftarrow (1 - \alpha_t)Q_t(s, a) + \alpha_t (r(s, a, s') + \delta V_t(s'))$;

$V_{t+1}(s) \leftarrow \max_a Q_t(s, a)$;

- ϵ -greedy: W.p. ϵ , act randomly, w.p. $(1 - \epsilon)$ act according to Q_t



Q-learning Properties

- Q-learning converges to optimal policy – even if agent acts sub-optimally!
- This is called **off-policy learning**
- There are some caveats
 - We have to explore enough
 - We have to eventually make the learning rate small enough
 - But we should not decrease it too quickly
 - Q-learning converges if $\sum_0^\infty \alpha_t = \infty$ and $\sum_0^\infty \alpha_t^2 < \infty$
 - Basically, in the limit, it does not matter how you select actions (!)



Outline

1. Introduction
2. Fictitious Play
3. Best-response Dynamics
4. No-regret Learning
5. Background: Single-agent Reinforcement Learning
- 6. Multi-agent Reinforcement Learning**



Independent Single-agent RL

- Setting: Two-player zero-sum games
- Naive idea: Agents **ignore** the existence of their opponent
- $Q_i^\pi(s, a_i)$: Value for i if both agents follow π starting from s and i plays a_i
- Learning dynamics: Agents deploy **independent Q-learning**
- **Good news**: **No-regret** property if opponent plays stationary policy
- **Bad news**: No convergence guarantee if both agents are learning (e.g., **self play**)!



Minimax-Q

- Littman² extended Q-learning algorithm to zero-sum stochastic games
- Main idea is to modify Q-function to consider actions of opponent

$$Q_{i,t+1}(s_t, a_t) = (1 - \alpha_t)Q_{i,t}(s_t, a_t) + \alpha_t (r_i(s_t, a_t) + \delta V_{i,t}(s_{t+1}))$$

- Since game is zero sum, we can have

$$V_{i,t}(s) = \max_{\pi_i} \min_{a_{-i}} Q_{i,t}(s, \pi_i, a_{-i})$$

²Littman, M. L. "Markov games as a framework for multi-agent reinforcement learning." 1994



Minimax-Q Algorithm

repeat *until convergence*

observe current state s ;

select action a_i and take it (e.g., via ϵ -greedy policy);

observe action profile a ;

observe next state s' and reward $r(s, a, s')$;

$Q_{i,t+1}(s, a) \leftarrow (1 - \alpha_t)Q_{i,t}(s, a) + \alpha_t (r(s, a) + \delta V_{i,t}(s'))$;

$\pi_i(s, \cdot) \leftarrow \operatorname{argmax}_{\pi'} \min_{a_{-i}} \sum_{a_i} \pi'(s, a_i) Q_{i,t}(s, a_i, a_{-i})$;

$V_{t+1}(s) \leftarrow \min_{a_{-i}} \sum_{a_i} \pi(s, a_i) Q_{i,t}(s, a_i, a_{-i})$;



Minimax-Q Algorithm: Discussion

- It guarantees agents payoff at least equal to that of their maxmin strategy
- In zero-sum games, minimax-Q converges to the value of the game in self play
- It no longer satisfies no-regret property
- If opponent plays sub-optimally, minimax-Q does not exploit it in most games



Nash-Q

- Hu and Wellman³ extended minimax-Q to general-sum games
- Algorithm is structurally identical to minimax-Q
- Extension requires that each agent maintains values for all other agents
- LP to find maxmin value is replaced with quadratic programming to find NE
- Nash-Q makes number of very limiting assumptions (e.g., uniqueness of NE)

³Hu, J, and Wellman, M. P. "Multiagent reinforcement learning: theoretical framework and an algorithm." 1998



Recall: Stochastic Games Model

- Focus on **stationary Markov strategies** (a mixed strategy per state)
- $\pi_i : S \mapsto \Delta(A_i)$ denotes (mixed) strategy of agent i at state s
- $\pi = (\pi_1, \dots, \pi_n)$ denotes strategy profile of all agents
- **Expected utility (value) function** of agent i is

$$v_i(s, \pi) := \mathbb{E}_{a_k \sim \pi(s_k)} \left[\sum_{k=0}^{\infty} \delta^k r_i(s_k, a_k) \mid s_0 = s \right]$$



Equilibrium Characterization

- Equilibrium **value function** is defined using **one-stage deviation** principle (**multi-agent** extension of **Bellman's equation**) as

$$v_i(s, \pi^*) = \max_{\pi_i} \mathbb{E}_{a \sim (\pi_i, \pi_{-i}^*(s))} \left[r_i(s, a) + \delta \sum_{s' \in S} p(s, a, s') v_i(s', \pi^*) \right]$$

- **Q-function** is defined as

$$Q_i(s, a, \pi^*) = r_i(s, a) + \delta \sum_{s' \in S} p(s, a, s') v_i(s', \pi^*)$$

- Recursion is then defined as

$$v_i(s, \pi^*) = \max_{\pi_i} \mathbb{E}_{a \sim (\pi_i, \pi_{-i}^*(s))} [Q_i(s, a, \pi^*)]$$



FP for Model-based Learning

- Consider learning dynamic that combines FP with value-function (or Q-function) iteration
- Agents form beliefs on opponent strategies (using empirical frequencies and assuming opponent uses stationary strategy)
- Agents also form beliefs about **equilibrium** value function, or Q-function
- Agents then choose best response action in **auxiliary game** given their beliefs (where payoffs are given by Q-function estimates)
- Key challenge is that payoffs or value functions in these auxiliary games are **non-stationary** (unlike repeated play of stage games)



FP for Model-based Learning: Model

- At time t , i 's belief on $-i$'s strategy is μ_i^t and on own Q-function is

$$Q_i^t := \mathbb{E}_{a_{-i} \sim \mu_{-i}^t(s)} [Q_i^t(s, a_i, a_{-i})]$$

- Agent i selects best response $a_i^t(s) \in \operatorname{argmax}_{a_i} Q_i^t(s, a_i, \mu_{-i}^t(s))$
- Agent i updates μ_i as

$$\mu_i^{t+1}(s) = (1 - \alpha_t)\mu_i^t(s) + \alpha_t a_{-i}^t(s)$$

- Agent i updates Q_i as

$$Q_i^{t+1}(s, a) = (1 - \beta_t)Q_i^t(s, a) + \beta_t \left(r_i(s, a) + \delta \sum_{s' \in S} p(s, a, s') v_i^t(s') \right)$$

where $v_i^t(s') = \max_{a_i} Q_i^t(s', a_i, \mu_{-i}^t(s))$



Two-timescale Learning Framework

- Beliefs on Q-functions are updated at slower rate than beliefs on opponent strategies
- This postulate agents' choices to be more dynamic than changes in their preferences
- Q-functions in auxiliary games can be viewed as slowly evolving agent preferences
- This enables weakening the dependence between evolving strategies and Q-functions



Convergence of Two-timescale Learning Framework

- If each state is visited **infinitely** many times
- And, if $\lim_{k \rightarrow \infty} \alpha_k = \lim_{k \rightarrow \infty} \beta_k = 0$ and $\sum_k \alpha_k = \sum_k \beta_k = \infty$
- And, if $\lim_{k \rightarrow \infty} \beta_k / \alpha_k = 0$ (two-timescale learning: $\beta_k \rightarrow 0$ faster than $\alpha_k \rightarrow 0$)
- Then Q and μ converge to NE value and strategy in zero-sum stochastic games
- They also converge to NE value for single-controller stochastic games



Acknowledgment

- This lecture is a slightly modified version of ones prepared by
 - Asu Ozdaglar [MIT 6.254]
 - Vincent Conitzer [Duke CPS 590.4]
 - Aaron Roth [UPenn NETS 412]
 - Dan Klein and Pieter Abbeel [UC Berkeley CS 188]

